

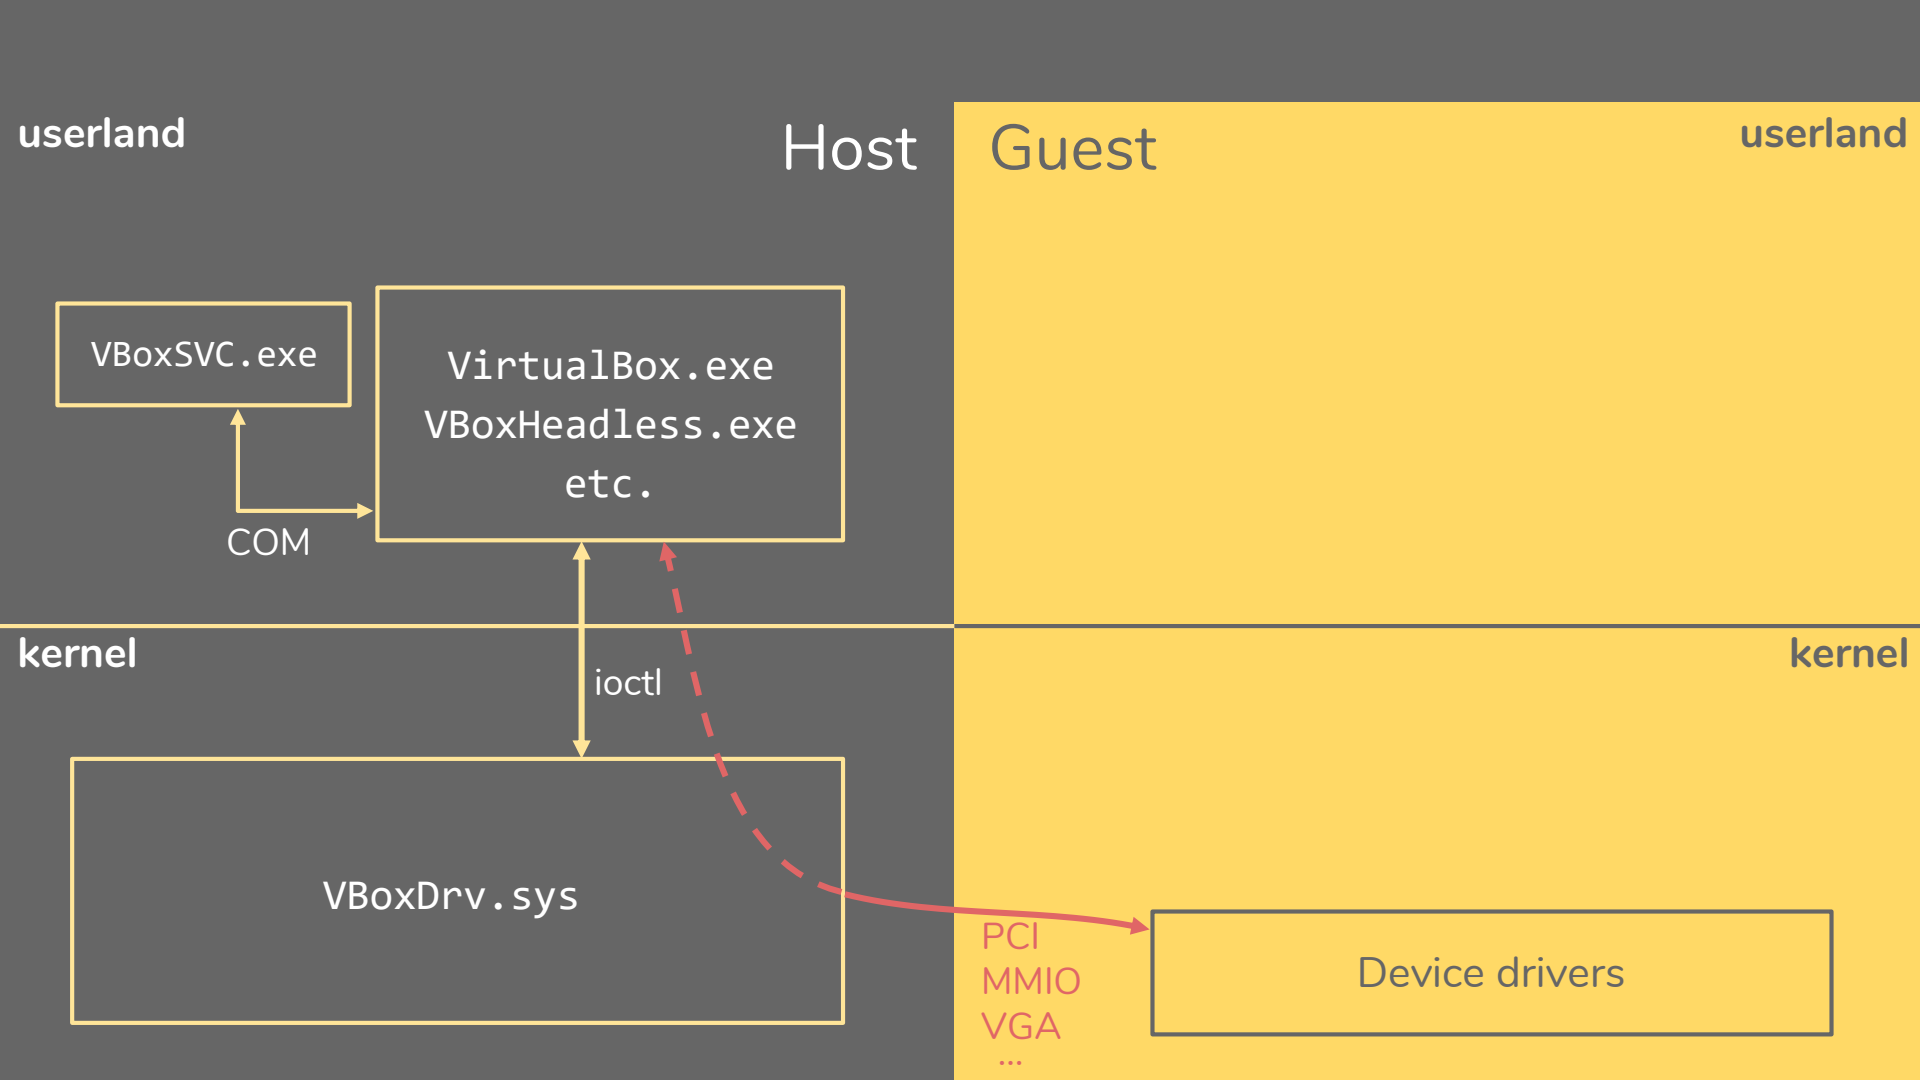
Thinking

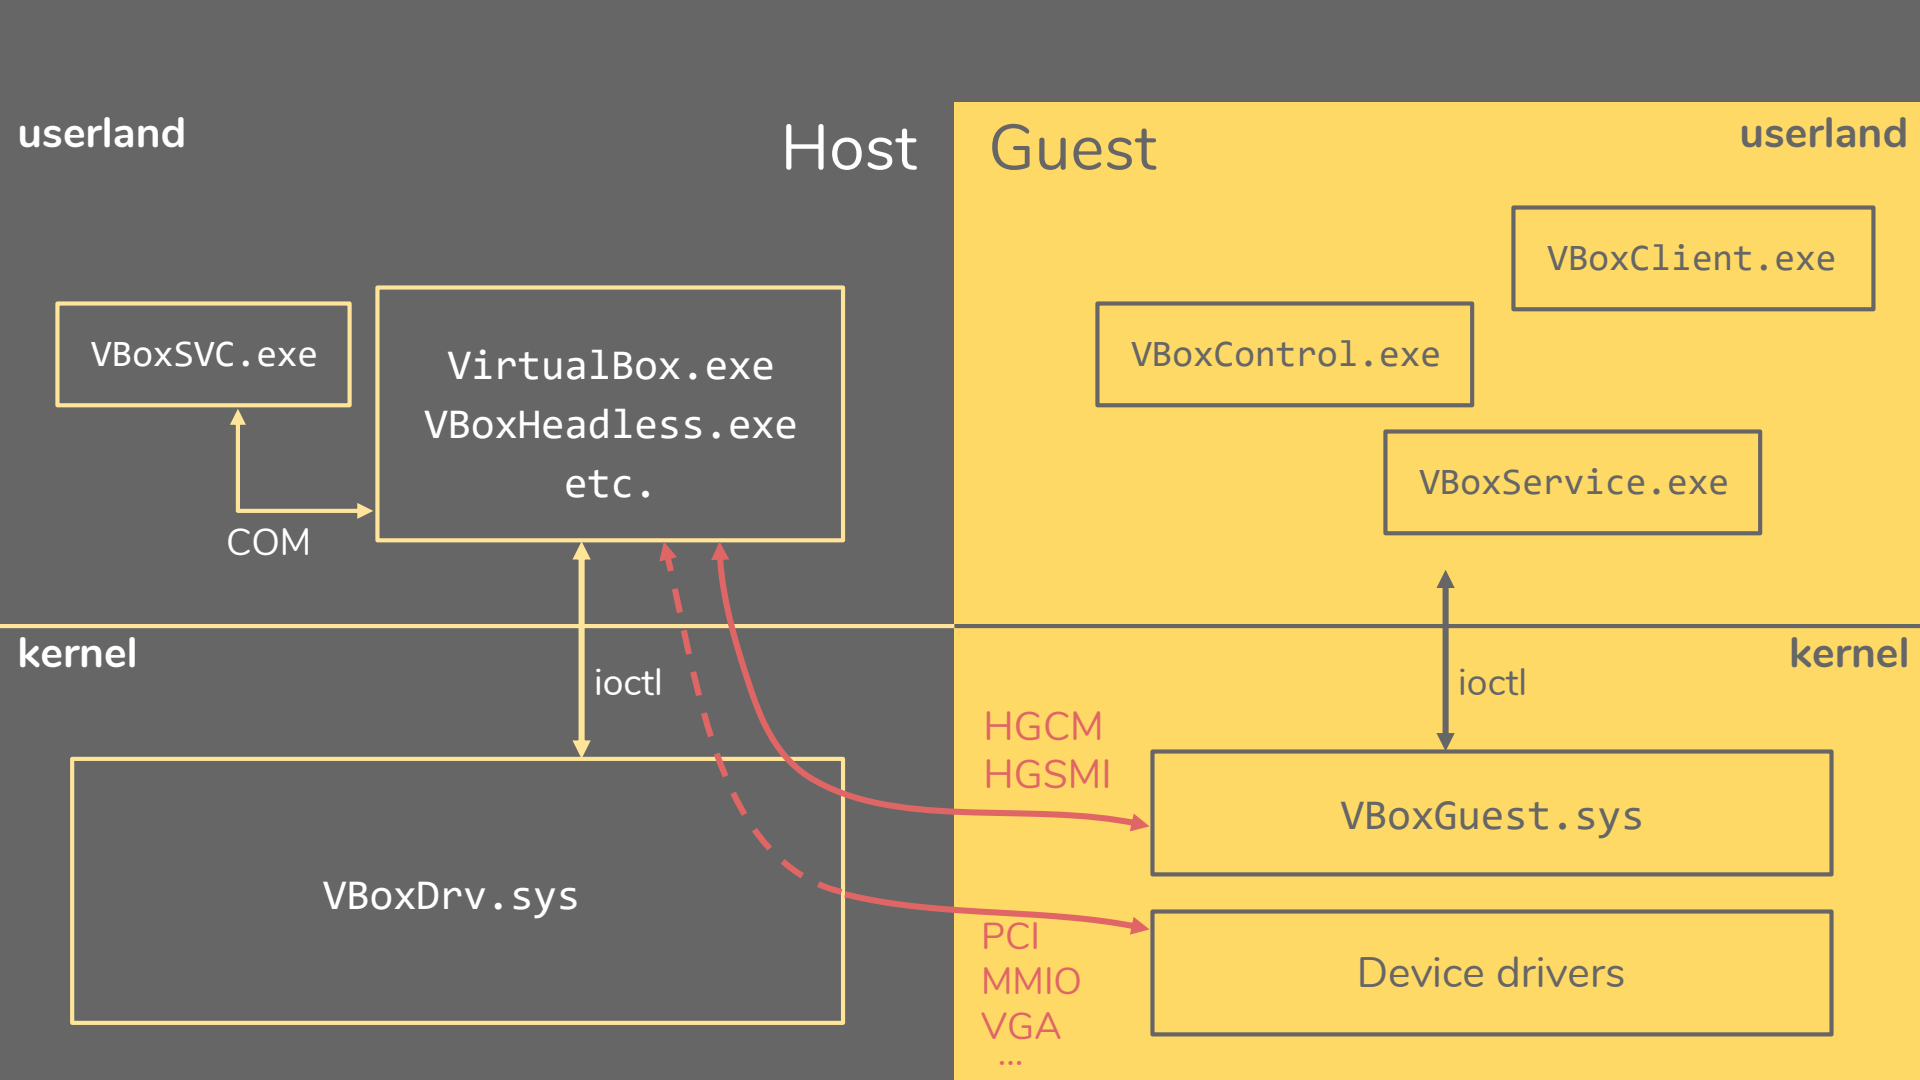
the (Virtual)Box

outside

Niklas Baumstark

 @_niklasb





Guest-to-host interfaces

- Hypervisor [/src/VBox/VMM](#)
 - Memory manager
 - x86 emulation
- Emulated devices [/src/VBox/Devices](#)
 - Audio
 - Networking
 - Graphics (VGA)
 - AHCI
 - ACPI
 - USB
 - Virtual Machine Monitor device
 - Paravirtualization interface (KVM/Hyper-V)
- HGCM services [/src/VBox/HostServices](#)
 - Shared OpenGL
 - Drag & Drop
 - Shared folders
 - Shared clipboard
- HGSMI services
 - VirtualBox Video Acceleration (VBVA)

Guest-to-host interfaces

- Hypervisor [/src/VBox/VMM](#)
 - Memory manager
 - x86 emulation
- Emulated devices [/src/VBox/Devices](#)
 - Audio
 - Networking
 - Graphics (VGA)
 - AHCI
 - ACPI
 - USB
 - Virtual Machine Monitor device
 - Paravirtualization interface (KVM/Hyper-V)
- HGCM services [/src/VBox/HostServices](#)
 - Shared OpenGL (20+ CVEs...)
 - Drag & Drop
 - Shared folders
 - Shared clipboard
- HGSMI services
 - VirtualBox Video Acceleration (VBVA)

Comparison to VMware Workstation

- Many features are disabled (= secure :) by default
 - 3D support
 - Drag & drop
 - Clipboard sharing
 - USB 2.0 & 3.0
- Some vectors do not exist
 - ThinPrint
- No userland RPC backdoor
- VirtualBox userland parts are privileged, privesc to host kernel is trivial

Host-Guest Communication Manager

- Simple RPC protocol, handled by the VMM PCI device
- Guest allocates request buffer of type
 - `VMMDevHGCMConnect` or
 - `VMMDevHGCMDisconnect` or
 - `VMMDevHGCMCall`
- Physical address of request is written to I/O port
- Call request specifies function ID and parameters
 - Integers
 - Buffers

HGCM - Services

- VBoxSharedClipboard
- VBoxDragAndDropSvc
- VBoxGuestPropSvc
- VBoxGuestControlSvc
- VBoxSharedFolders
- VBoxSharedCrOpenGL

HGCM - Services

- VBoxSharedClipboard
- VBoxDragAndDropSvc
- **VBoxGuestPropSvc**
- **VBoxGuestControlSvc**
- VBoxSharedFolders
- VBoxSharedCrOpenGL

HGCM - Example

```
C:\Users\niklas>VBoxControl.exe guestproperty set foo bar
Oracle VM VirtualBox Guest Additions Command Line Management Interface Version 5.2.8
(C) 2008-2018 Oracle Corporation
All rights reserved.

C:\Users\niklas>VBoxControl.exe guestproperty get foo
Oracle VM VirtualBox Guest Additions Command Line Management Interface Version 5.2.8
(C) 2008-2018 Oracle Corporation
All rights reserved.


Value: bar

C:\Users\niklas>
```

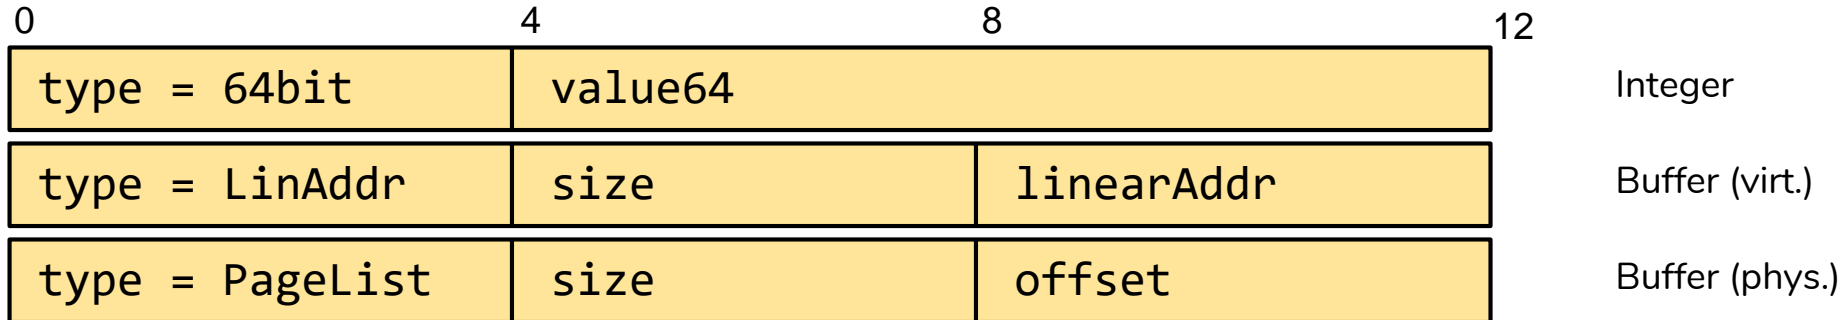
```
VMMDevHGCMConnect("VBoxGuestPropSvc") = 42
VMMDevHGCMCall(42, SET_PROP, "foo", "bar") = VERR_SUCCESS
VMMDevHGCMCall(42, GET_PROP, "foo", <result buffer>, ...) = VERR_SUCCESS
```

HGCM - Call request

VMMDevHGCMCall




HGCMFunctionParameter32 (12 bytes)

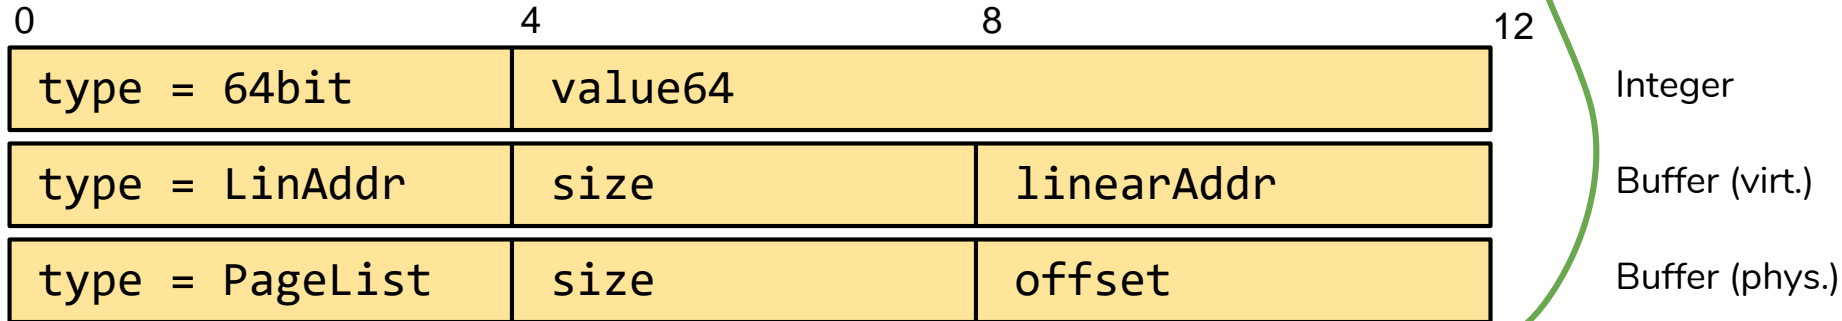


HGCM - Call request

VMMDevHGCMCall




HGCMFunctionParameter32 (12 bytes)



HGCM - Call handling (< 5.2.10)

1. Copy `VMMDevHGCMCall` to host heap
2. Allocate `VBOXHGCMCMD` large enough to hold copy of parameters (host params) & buffer data
3. Copy buffer data from the guest into the `VBOXHGCMCMD`



Bug #1: Double fetch on buffer write-back



- Most HGCM functions return data
- Implemented by writing back `VBOXHGCMCMD` buffers to guest memory
- `hgcmCompletedWorker` re-fetches the request to determine sizes
- Disclose heap memory by increasing the size during dispatch!

```
case VMMDevHGCMParamType_LinAddr: {
    /* Copy buffer back to guest memory. */
    uint32_t size = pGuestParam->u.Pointer.size;
    ...
    /* Use the saved page list to write data back to the guest RAM. */
    rc = vmmdevHGCMWriteLinPtr (... ,
                                pHostParam->u.pointer.addr,
                                size, ...);
}
```


Bug #2: Heap out-of-bounds

read



- `VMMDevHGCMCall` is copied from guest to the host heap
 - `size` bytes are copied
 - No check that `size` is large enough to hold all parameters
 - Later: OOB read access on the heap
- Looks harmless, because the guest fully controls the object anyways


`VMMDevHGCMCall`



Bug #2: Heap out-of-bounds **double** read




- `VMMDevHGCMCall` is copied from guest to the host heap
 - `size` bytes are copied
 - No check that `size` is large enough to hold all parameters
 - Later: OOB read access on the heap
- Looks harmless, because the guest fully controls the object anyways
- But: **parameters are accessed twice!** TOCTOU issue?



Bug #2: Heap out-of-bounds **double** read




- **Pass 1:** `cbCmdSize` variable sums up `size` values for buffer params
- Allocate `VBOXHGCMCMD` including space for `cbCmdSize` bytes of data
- **Pass 2:** Copy data from guest
 - Due to OOB access we can change `size` values concurrently



<code>type = LinAddr</code>	<code>size</code>	<code>linearAddr</code>
-----------------------------	-------------------	-------------------------

<code>type = PageList</code>	<code>size</code>	<code>offset</code>
------------------------------	-------------------	---------------------

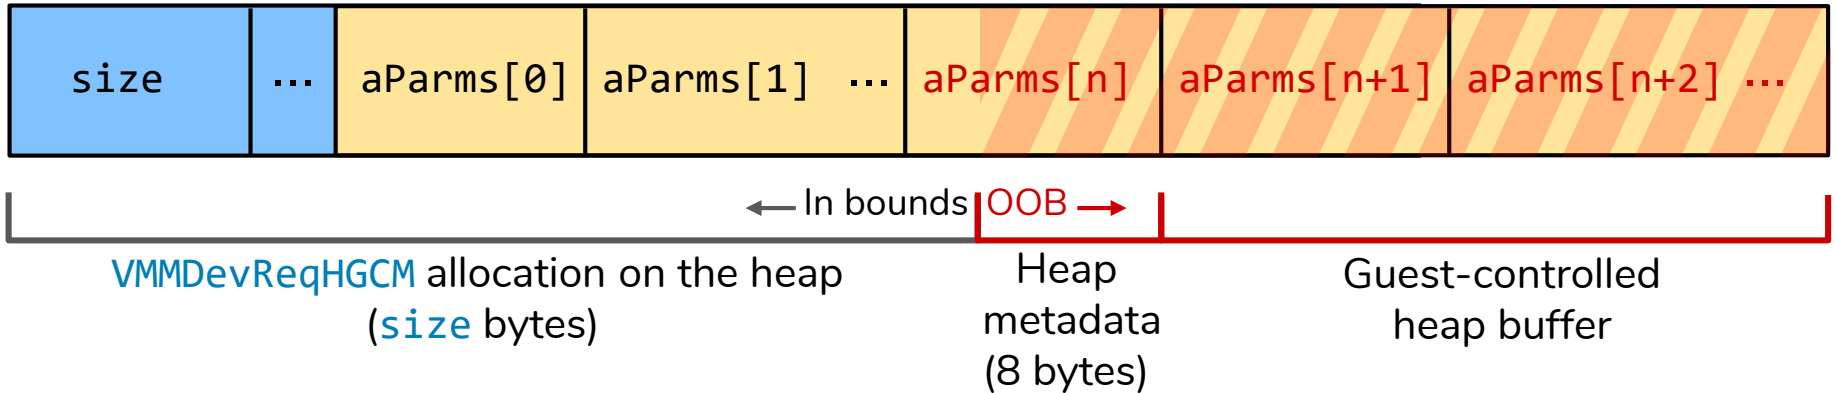
Bug #2: Turn into OOB write



Params $n + 1$ and upwards are in a different heap chunk,
we can race them between the two passes

=> Requires two vCPUs

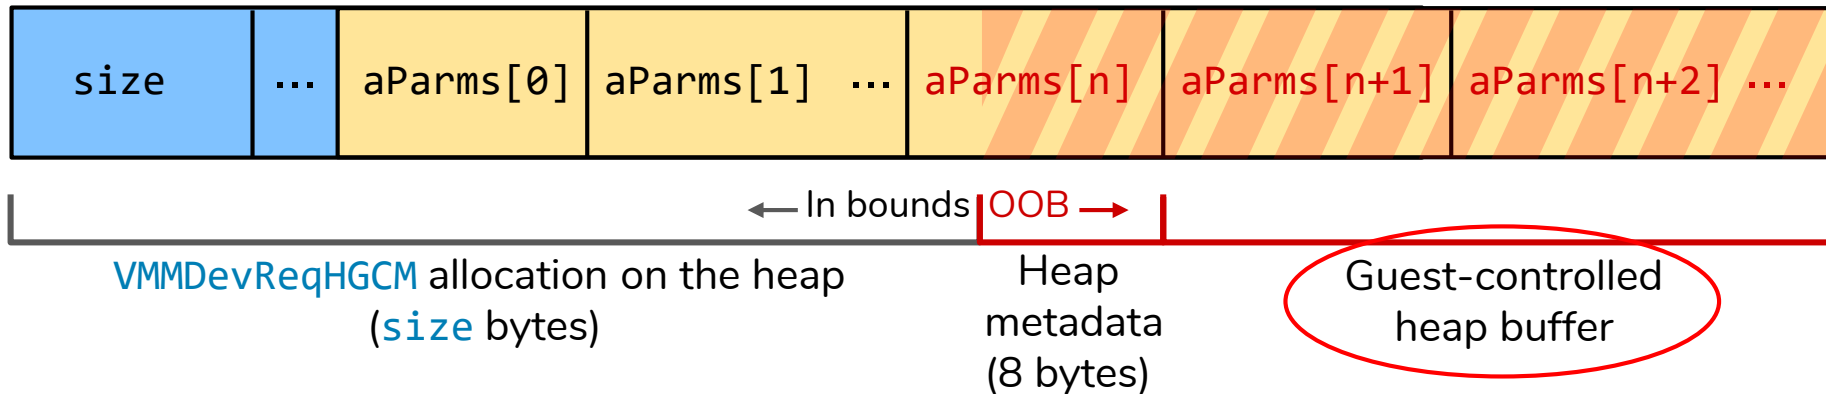
Bug #2: Turn into OOB write



Challenge 1: Find an object on the heap that we can write repeatedly

Challenge 2: Incorporate heap metadata into the request

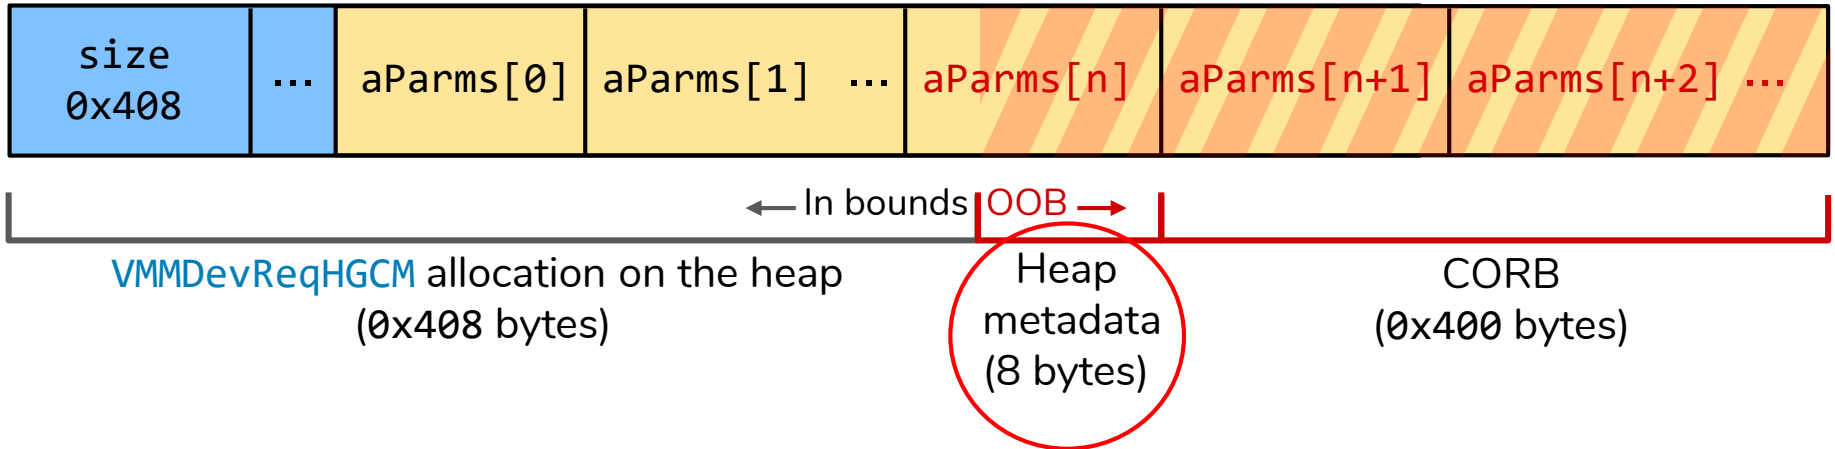
Bug #2: Turn into OOB write



Find an object on the heap that we can write repeatedly

- Intel HD audio device command output ring buffer (CORB)
 - 0x400 bytes, can be re-allocated at will

Bug #2: Turn into OOB write



Incorporate heap metadata into request

- `CORB` size `0x400` => LFH bucket size `0x410` (incl. 8 bytes metadata)
- `VMMDevReqHGCM` with 83 parameters has size `0x410`
 - Last 8 bytes of 83rd parameter are uncontrolled heap metadata
 - This is ok for integer parameters!

Bug #2: Make it an OOB write

One thread constantly flips a `PageList` parameter size in CORB

```
uint32_t size = pGuestParm->u.PageList.size; // <- fully controlled!  
...  
// This will happily read less than size bytes, if page list is smaller  
rc = vmmdevHGCMPageListRead(pThis->pDevIns, pcBuf, size, pPageListInfo);  
...  
pcBuf += size; // <- will be used as destination for the next parameter
```

Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410



Segment heap (> 8k)



Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410




Segment heap (> 8k)



Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410




Segment heap (> 8k)



Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410




Segment heap (> 8k)



Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410




Segment heap (> 8k)



Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410




Segment heap (> 8k)




Exploit

- Powerful, heap-based **relative** read and write primitives
- **VBOXHGCMCMD** is variable-size, we can put it on a “nice” heap
 - Allocator fully predictable, with help of bug #1
- Can already leak some vtable pointers from **VBoxC.dll** with bug #1
- Next: want to corrupt
 - a pointer that is read from (for full ASLR break)
 - a function or vtable pointer (for control flow hijack)

LFH bucket 0x410

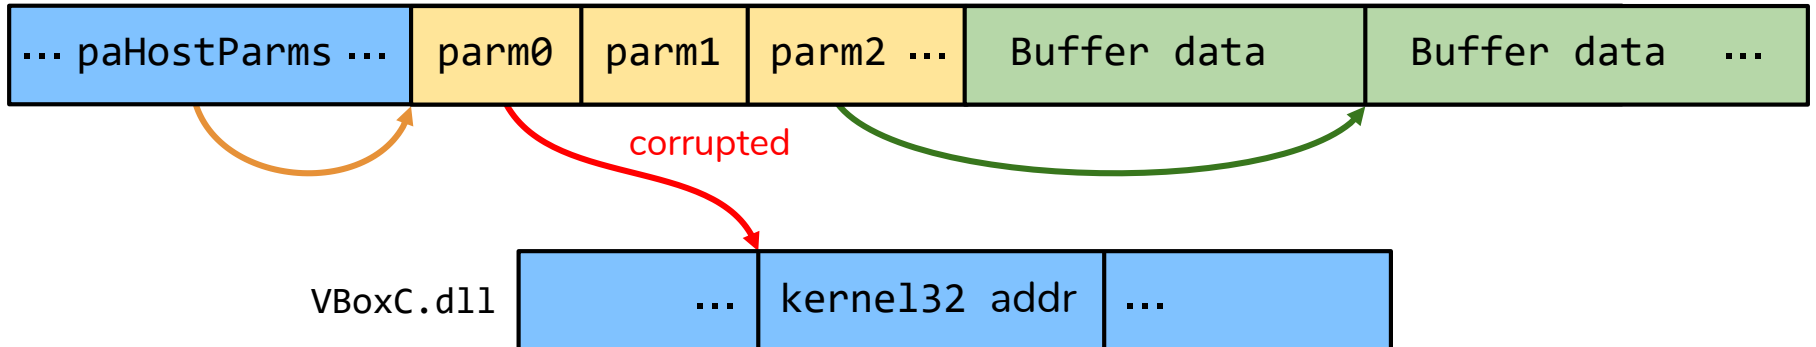


Segment heap (> 8k)



Exploit - Absolute read

- **VBOXHGCMCMD** is an interesting data structure with pointers
 - Not all HGCM calls return immediately!
 - Send **GET_NOTIFICATION** to guest properties service
 - It returns when a property is set that matches the given pattern
 - This will cause a writeback using pointers from **VBOXHGCMCMD**
 - Used to leak **kernel32** and **ntdll** base addresses



Exploit - Nail in the coffin

- A `HGCMMsgCall` object is allocated for each HGCM call
- Unlike `VBOXHGCMCMD`, it has a constant size of `0x98` => LFH heap
- Contains a pointer to itself
 - use a small spray and bug #1 to find it
- We corrupt the `pHGCMPort` field

Segment heap (> 8k)




LFH bucket 0xa0



Exploit - Nail in the coffin

- A `HGCMMsgCall` object is allocated for each HGCM call
- Unlike `VBOXHGCMCMD`, it has a constant size of `0x98` => LFH heap
- Contains a pointer to itself
 - use a small spray and bug #1 to find it
- We corrupt the `pHGCMPort` field



Exploit - Nail in the coffin

- A `HGCMMsgCall` object is allocated for each HGCM call
- Unlike `VBOXHGCMCMD`, it has a constant size of `0x98` => LFH heap
- Contains a pointer to itself
 - use a small spray and bug #1 to find it
- We corrupt the `pHGCMPort` field

```
static DECLCALLBACK(void) hgcmMsgCompletionCallback(int32_t result,
                                                    HGCMMsgCore *pMsgCore) {
    /* Call the VMMDev port interface to issue IRQ notification. */
    HGCMMsgHeader *pMsgHdr = (HGCMMsgHeader *)pMsgCore;
    ...
    if (pMsgHdr->pHGCMPort && !g_fResetting) {
        pMsgHdr->pHGCMPort->pfnCompleted(pMsgHdr->pHGCMPort, ...);
    }
}
```

Oracle VM VirtualBox Manager

File Machine Help

New Settings Discard Show

Snapshots Machine Tools Global Tools

Take Delete Restore Properties Clone

Name	Taken
<ul style="list-style-type: none"> installe <ul style="list-style-type: none"> pwn <ul style="list-style-type: none"> Current State (chang...) 	22 Aug 2018 16:45 (22 days ago) 23 Aug 2018 09:09 (21 days ago)

Recycle Bin

Google Chrome

Firefox

Oracle VM VirtualBox

VMware Workstati...

procexp64

Recycle Bin

download chrome

rootshell

From `VirtualBox.exe` to host kernel

- `VirtualBox.exe` is privileged, since it has access to `VBoxDrv` IOCTLs
- `SUP_IOCTL_LDR_{OPEN,LOAD}` load PE file as kernel plugin
 - Verifies driver signature
- `SUP_IOCTL_CALL_SERVICE` calls into a plugin
 - Full control over 4th argument
 - => RIP control via `jmp r9`
- `SUP_IOCTL_PAGE_{ALLOC_EX,MAP_KERNEL,PROTECT}`
 - Map RWX code in the kernel

Because why not



CVE-2018-2698

- HGSMI = **H**ost-**G**uest **S**hared **M**emory **I**nterface
- Guest allocates request buffer in video RAM, notifies VGA device
- Used for VBVA subsystem (**V**irtual**B**ox **V**ideo **A**cceleration)
- **VBVA_VDMA_CMD** is used for video DMA commands:
 - **VBOXVDMACMD_TYPE_DMA_PRESENT_BLT**
 - **VBOXVDMACMD_TYPE_DMA_BPB_TRANSFER**

```
int rc = vboxVDMACmdExecBltPerform(pVdma, pvRam + pBlt->offDst, pvRam + pBlt->offSrc,
    &pBlt->dstDesc, &pBlt->srcDesc,
    pDstRectl,
    pSrcRectl);

...
static int vboxVDMACmdExecBltPerform(PVBOXVDMAHOST pVdma,
    uint8_t *pvDstSurf, const uint8_t *pvSrcSurf,
    const PVBOXVDMA_SURF_DESC pDstDesc, const PVBOXVDMA_SURF_DESC pSrcDesc,
    const VBOXVDMA_RECTL * pDstRectl, const VBOXVDMA_RECTL * pSrcRectl)
{
    ...
    if (pDstDesc->width == pDstRectl->width && pSrcDesc->width == pSrcRectl->width
        && pSrcDesc->width == pDstDesc->width) {
        ...
        uint32_t cbOff = pDstDesc->pitch * pDstRectl->top;
        uint32_t cbSize = pDstDesc->pitch * pDstRectl->height;
        memcpy(pvDstSurf + cbOff, pvSrcSurf + cbOff, cbSize);
    }
}
```

= guest-controlled

```
int rc = vboxVDMACmdExecBltPerform(pVdma, pvRam + pBlt->offDst, pvRam + pBlt->offSrc,
    &pBlt->dstDesc, &pBlt->srcDesc,
    pDstRect1,
    pSrcRect1);

...

static int vboxVDMACmdExecBltPerform(PVBOXVDMAHOST pVdma,
    uint8_t *pvDstSurf, const uint8_t *pvSrcSurf,
    const PVBOXVDMA_SURF_DESC pDstDesc, const PVBOXVDMA_SURF_DESC pSrcDesc,
    const VBOXVDMA_RECTL * pDstRect1, const VBOXVDMA_RECTL * pSrcRect1)
{
    ...
    if (pDstDesc->width == pDstRect1->width && pSrcDesc->width == pSrcRect1->width
        && pSrcDesc->width == pDstDesc->width) {
        ...
        uint32_t cbOff = pDstDesc->pitch * pDstRect1->top;
        uint32_t cbSize = pDstDesc->pitch * pDstRect1->height;
        memcpy(pvDstSurf + cbOff, pvSrcSurf + cbOff, cbSize);
    }
}
```

VirtualBox host debugging

- Cannot attach to `VirtualBox.exe` due to *process hardening*
- Exploit dev on Windows: non-hardened debug build
 - Get ready for a nostalgic experience with VS 2010
 - Ideally have a friend do it for you
- Debugging the official Windows build:
 - Run VirtualBox inside VMware Workstation (enable “Virtualize Intel VT-x”)
 - Use a kernel debugger with `!gflag +soe` and `!process`
- Bug hunting & PoCs are much easier on Linux host + guest
 - Configure guest VM according to target

Dig deeper

Advisories for presented bugs <https://www.zerodayinitiative.com/advisories/ZDI-18-782/> <https://www.zerodayinitiative.com/advisories/ZDI-18-783/> <https://blogs.securiteam.com/index.php/archives/3649>

Bugs in E1000 network card, NAT & virtio-net (2017) https://github.com/fundacion-sadosky/vbox_cve_2017_10235
<https://bugs.chromium.org/p/project-zero/issues/detail?id=1086> <https://bugs.chromium.org/p/project-zero/issues/detail?id=1136>

VDMA exploit and host-/guest-based privilege escalations (2018) <https://www.youtube.com/watch?v=fFaWE3jt7qU>
https://github.com/phoenix/files/blob/master/slides/unboxing_your_virtualboxes.pdf

VBVA double fetch (2018) <https://www.voidsecurity.in/2018/08/from-compiler-optimization-to-code.html>

Windows process hardening <https://googleprojectzero.blogspot.com/2017/08/bypassing-virtualbox-process-hardening.html>

VirtualBox 3D hacks <https://www.coresecurity.com/corelabs-research/publications/breaking-out-virtualbox-through-3d-acceleration>
<https://phoenix.re/2018-07-27/better-slow-than-sorry> <https://github.com/niklasb/3dpwn>
<https://www.thezdi.com/blog/2018/8/28/virtualbox-3d-acceleration-an-accelerated-attack-surface>

Simple Python HGCM client library: <https://github.com/niklasb/3dpwn/blob/master/lib/hgcm.py>

VMware Workstation vulnerabilities & exploitation <https://keenlab.tencent.com/en/2018/04/23/A-bunch-of-Red-Pills-VMware-Escapes/>
<https://www.thezdi.com/blog/2018/3/1/vmware-exploitation-through-uninitialized-buffers>
https://comsecuris.com/blog/posts/vmware_vgpu_shader_vulnerabilities/ <https://www.blackhat.com/docs/eu-17/materials/eu-17-Mandal-The-Great-Escapes-Of-Vmware-A-Retrospective-Case-Study-Of-Vmware-G2H-Escape-Vulnerabilities.pdf> and many more

DSEFix (exploits old VBoxDrv version to disable Driver Signature Enforcement on Windows 10 \leq RS4) <https://github.com/hfiref0x/DSEFix>
fwexp1: awesome framework for low-level I/O hacking (with ridiculous RWEverything driver) <https://github.com/Cr4sh/fwexp1>